

Introductory Econometrics with Recitation: TA Session Week 7

Danny Po-Hsien Kang (康柏賢)

April 7, 2026

Today's agenda

- 1 Random Number Generator
 - `set.seed()`
 - `runif()`
 - `rnorm()`
 - `rbinom()`
 - `rpois()`
 - `rexp()`
- 2 Data Wrangling: Combine Dataset
 - `rbind()`
 - `cbind()`
- 3 Standardization
 - `scale()`
- 4 Multiple Testing
 - `linearHypothesis()`

Today's Dataset

- Today we will use the following CSV file from the OpenIntro Website:
 - `ncbirths.csv`: A 2004 dataset from North Carolina on births, including mothers' habits and practices, with 1,000 cases.
- Please import the following datasets into RStudio.

```
birth <- read.csv("data/ncbirths.csv")
```

Set Random Seed: `set.seed()`

- `set.seed()` sets R's random-number generator (RNG) state
 - Random sampling functions: `runif()`, `rnorm()`, `sample()`
 - Without setting seeds, we will obtain different outcomes every time we use these functions.
 - Key to **reproducibility**: It guarantees that the code produces the same sequence every time you rerun.
- **Syntax:**

```
set.seed(seed)
```

- **Example:**

```
set.seed(20260407)
```

Generate Uniform Random Samples: `runif()`




Figure: Random samples drawn from $\text{Uniform}[0,1]$ in different sizes

Generate Uniform Random Samples: runif()

- **Syntax:**

```
num <- runif(n = ..., min = ..., max = ...)
```

- **Example:**

```
rdnum <- runif(10) # min = 0, max = 1  
rdnum_2 <- runif(n = 10, min = 0, max = 100)
```

- **Output**

```
print(rdnum)  
[1] 0.7996254 0.7107656 0.4796840 0.6867655 0.7798512  
[6] 0.6304974 0.6980979 0.8342457 0.5835580 0.9687111  
  
print(rdnum_2)  
[1] 54.44931 48.50616 99.31460 60.44028 64.11838  
[6] 65.01593 38.41537 46.73956 28.80531 13.15521
```

Generate Normal Random Samples: `rnorm()`




Figure: Random samples drawn from $\mathcal{N}(0, 1)$ in different sizes

Generate Normal Random Samples: rnorm()

- **Syntax:**

```
num <- rnorm(n = ..., mean = ..., sd = ...)
```

- **Example:**

```
rdnum <- rnorm(10) # mean = 0, sd = 1  
rdnum_2 <- rnorm(n = 10, mean = 60, sd = 15)
```

- **Output**

```
print(rdnum)  
[1]  0.8360853 -0.6300108 -0.0017627  2.2970599 -0.8165532  
[6] -2.0158603  0.2546558  0.1085037  0.0042885  0.4861076  
  
print(rdnum_2)  
[1] 66.06636 71.63407 75.30722 60.50602 46.48625  
[6] 52.42269 67.73190 41.56804 72.80670 61.29849
```

Generate Binomial Random Samples: `rbinom()`




Figure: Random samples drawn from $\text{Bin}(20, 0.2)$ in different sizes

Generate Binomial Random Samples: rbinom()

- **Syntax:**

```
a <- rbinom(n = ..., size = ..., prob = ...)
```

- **Example:**

```
rdnum <- rbinom(n = 30, size = 20, prob = .5)
```

- **Output**

```
print(rdnum)
[1] 10 11  9 11 13  8 11 11  8 12
[11]  8  7 13 10  8 13 12 11 11 10
[21] 10  7 10 10 10 15  8  7  8 11

print(mean(rdnum))
[1] 10.1
```

Generate Poisson Random Samples: `rpois()`




Figure: Random samples with size 1,000 drawn from $\text{Poisson}(\lambda = 5)$

Generate Poisson Random Samples: `rpois()`

- **Syntax:**

```
num <- rpois(n = ..., lambda = ...)
```

- **Example:**

```
rdnum_pois <- rpois(20, lambda = 5)
```

- **Output:**

```
> print(rdnum_pois)
[1] 3 4 5 1 3 7 8 4 5 10 8 6 3 11 1 4 4 3 7 4
```

Generate Exponential Random Samples: `rexp()`




Figure: Random samples with size 1,000 drawn from $\text{Exp}(\lambda = 5)$

Generate Exponential Random Samples: rexp()

- **Syntax:**

```
num <- rexp(n = ..., rate = ...)
```

- **Example:**

```
rdnum_exp <- rexp(20, rate = 5)
```

- **Output:**

```
> print(rdnum_exp)
[1] 0.032601 0.053080 0.219614 0.615566 0.047792
[6] 0.112196 0.194019 0.119259 0.960832 0.462925
[11] 0.113932 0.305901 0.493771 0.236433 0.018397
[16] 0.111395 0.292260 0.293656 0.203214 0.138105
```

Two Ways to Combine Your Dataset




Figure: Combine dataset by rows or columns (Gaston Sanchez, 2024)

Combine Dataset by Rows: rbind()

- **Example:**

```
## first, let's draw two different sets of random sample
rdnum_pois_20 <- rpois(n = 20, lambda = 5)
rdnum_pois_10000 <- rpois(n = 10000, lambda = 5)

## using the two sets of random sample, build two DataFrames
df_20 <- data.frame(
  size = 20,
  values = rdnum_pois_20
)
df_10000 <- data.frame(
  size = 10000,
  values = rdnum_pois_10000
)

## combine the two DataFrames together, and draw the histogram
df <- rbind(df_20, df_10000)
```

Combine Dataset by Rows: rbind()

- **Example:**

```
ggplot(df, aes(x = values)) +  
  geom_histogram() +  
  facet_wrap(~size, scales = "free_y")
```

Combine Dataset by Rows: rbind()




Figure: Histogram of two samples in different sample sizes

Combine Dataset by Columns: cbind()

- **Example:**

```
## first, let's draw two different sets of random sample
rdnum_pois_param5 <- rpois(n = 1000, lambda = 5)
rdnum_pois_param20 <- rpois(n = 1000, lambda = 20)

## using the two sets of random sample, build two DataFrames re
df_param5 <- data.frame(
  values_param5 = rdnum_pois_param5
)
df_param20 <- data.frame(
  values_param20 = rdnum_pois_param20
)
```

Combine Dataset by Columns: cbind()

- **Example:**

```
df <- cbind(df_param5, df_param20)

ggplot(df) +
  geom_histogram(mapping = aes(x = values_param5,
                              fill = "lambda = 5"),
                 alpha = 0.7
  ) +
  geom_histogram(mapping = aes(x = values_param20,
                              fill = "lambda = 20"),
                 alpha = 0.7
  )
```

Combine Dataset by Columns: `cbind()`




Figure: Histogram of two samples in different distribution parameters

Standardize Variables: scale()

- **Syntax:**

```
new_var <- scale(x = ..., center = TRUE, scale = TRUE)
```

- **Example:**

```
score <- c(70, 80, 75, 95)  
scale(score)
```

- **Output:**

```
      [,1]  
[1,] -0.9258201  
[2,]  0.0000000  
[3,] -0.4629100  
[4,]  1.3887301  
attr(,"scaled:center")  
[1] 80  
attr(,"scaled:scale")  
[1] 10.80123
```

Standardize Variables: scale()

- If center = TRUE, then scale() subtracts the mean: $x_i - \bar{x}$
- If scale = TRUE, then scale() divides by the standard deviation $\frac{x_i}{s_x}$
- So when center = TRUE and scale = TRUE, $z_i = \frac{x_i - \bar{x}}{s_x}$
- **Example:**

```
scale(score, center = FALSE, scale = TRUE)
```

- **Output:**

```
      [,1]  
[1,] 0.7526447  
[2,] 0.8601653  
[3,] 0.8064050  
[4,] 1.0214463  
attr(,"scaled:scale")  
[1] 93.00538
```

Multiple Testing: linearHypothesis()

- Install and load the package:

```
install.packages("car")  
library(car)
```

- **car**: Short for Companion to Applied Regression.
- We will use this package to conduct multiple hypothesis testing.

Multiple Testing: linearHypothesis()

- Consider the following model:

$$\text{Weight}_i = \beta_0 + \beta_1 \text{Week}_i + \beta_2 \text{Gender}_i + \beta_3 \text{Visit}_i + \varepsilon_i$$

- We use the `lm_robust()` function to estimate the model.

```
lm_model_multiple_robust <- lm_robust(  
  data = birth,  
  weight ~ weeks + gender + visits  
)
```

Multiple Testing: `linearHypothesis()`

• Expected output:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-6.208828047	0.532881547	-11.651423	1.745020e-29
weeks	0.339617332	0.013969817	24.310794	1.162998e-102
gendermale	0.365472826	0.070476825	5.185716	2.610867e-07
visits	0.009322152	0.009229467	1.010042	3.127228e-01

- Single hypothesis testing results for each coefficient can be directly checked from the regression output:
 - standard error
 - t-value
 - p-value
- To test multiple coefficients simultaneously (joint hypothesis), we need to use the `linearHypothesis()` function.

Multiple Testing: linearHypothesis()

- **Syntax:**

```
library(car) # Companion to Applied Regression

linearHypothesis(
  Linear_Model,
  c("Hypothesis 1", "Hypothesis 2", ...),
  test = "F"
)
```

- **Example:**

```
linearHypothesis(
  lm_model_multiple_robust,
  c("weeks = 0", "visits = 0"), # or c("weeks = visits")
  test = "F", # or Chisq
)
```

Multiple Regression: Multiple Testing

• Output:

```
Linear hypothesis test:
```

```
weeks = 0
```

```
visits = 0
```

```
Model 1: restricted model
```

```
Model 2: weight ~ weeks + gender + visits
```

	Res.Df	Df	F	Pr(>F)
1	988			
2	986	2	309.37	< 2.2e-16 ***

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1
```